

Example Analysis using BiocMAP Output Objects

***Nicholas J. Eagles*^{*1} and *Leonardo Collado-Torres*^{†1,2}**

¹Lieber Institute for Brain Development, Johns Hopkins Medical Campus

²Center for Computational Biology, Johns Hopkins University

*nickeagles77@gmail.com †lcolladotor@gmail.com

1 March 2022

Contents

1	Add Experiment Metadata to BiocMAP Outputs	2
2	Exploratory Plots	3
2.1	Bisulfite-Conversion Efficiency by Cell Population	3
2.2	Relationship between Methylation Fractions across Cytosine Context, by Cell Population	4
2.3	Fraction of Covered Cytosines by Cell Population	4
2.4	Distribution of Methylation Fractions across Cytosines by Cell Population	5
3	Bibliography	13

Example Analysis using BiocMAP Output Objects

BiocMAP is a computational pipeline for bisulfite-sequencing data. Starting from the raw FASTQ-based sequencer outputs, BiocMAP ultimately produces two `bsseq` R objects, containing methylation proportions and coverage information for all genomic cytosines in CpG and CpH methylation context. Various metrics, such as bisulfite-conversion rate, number of reads trimmed, or concordant alignment rate, are collected and become columns in the `colData` of the `bsseq` objects. The below analysis will demonstrate how to explore both methylation information and quality metrics in the BiocMAP output objects. Background understanding of the `bsseq` package, and its parent `SummarizedExperiment` are helpful but not required here.

1 Add Experiment Metadata to BiocMAP Outputs

The `bsseq` output objects from BiocMAP contain methylation and coverage info for our samples in the dataset. However, we're interested in exploring how this information relates back to sample metadata and phenotype information, present in an external file. Our first step will therefore be to load the BiocMAP output objects into memory, and manually attach the additional sample metadata to each object.

```
# Load required R packages
library("bsseq")
library("HDF5Array")
library("ggplot2")
library("GGally")
library("tidyr")
```

```
# Path to the sample metadata and BiocMAP outputs. The outputs are too large
# to host in this repository, so we reference local paths here
meta_file <- file.path(
  "/dcl02/lieber/ajaffe/FlowRNA_RNAseq/WGBS",
  "FlowRNA_WGBS_Sample_Information_with_Phenotype_Info.csv"
)
out_dir <- file.path(
  "/dcs04/lieber/lcolladotor/flowRNA_LIBD001/flowRNA_WGBS/processed-data",
  "03-BiocMAP/BiocMAP-output"
)
```

```
# Load the 'CpG'-context object
bs_cpg <- loadHDF5SummarizedExperiment(
  file.path(out_dir, "BSobjects", "objects", "combined"),
  prefix = "CpG"
)

# Load the 'CpH'-context object. Note: this requires quite a bit of memory
# (~23GB) even though the assays are disk-backed!
bs_cph <- loadHDF5SummarizedExperiment(
  file.path(out_dir, "BSobjects", "objects", "combined"),
  prefix = "CpH"
)

# Read in experiment-specific metadata and ensure sample ID orders match
meta <- read.csv(meta_file)
```

Example Analysis using BiocMAP Output Objects

```
meta <- meta[match(colnames(bs_cpg), meta$LIBD.), ]

# Add this metadata to the Bioconductor objects
colData(bs_cpg) <- cbind(colData(bs_cpg), meta)
colData(bs_cph) <- cbind(colData(bs_cph), meta)

# Keep a copy of the metadata as a data frame, for easy plotting
meta_df <- data.frame(colData(bs_cpg))
```

2 Exploratory Plots

2.1 Bisulfite-Conversion Efficiency by Cell Population

This experiment used spike-ins of the lambda bacteriophage genome, which were quantified via BiocMAP to infer bisulfite-conversion rate. Successful bisulfite conversion is a pre-requisite for accurate methylation calls, so we'd like to see both that values (interpreted as percentages) are close to 100, and that values are not significantly different by sample (or by sample-related variables like cell population). We'll explore this visually below.

```
ggplot(meta_df, aes(x = Cell.Population, y = lambda_bs_conv_eff)) +
  geom_boxplot(outlier.shape = NA) +
  geom_point() +
  labs(x = "Cell Population", y = "Bisulfite Conversion Rate (%)") +
  theme_bw(base_size = 20)
```

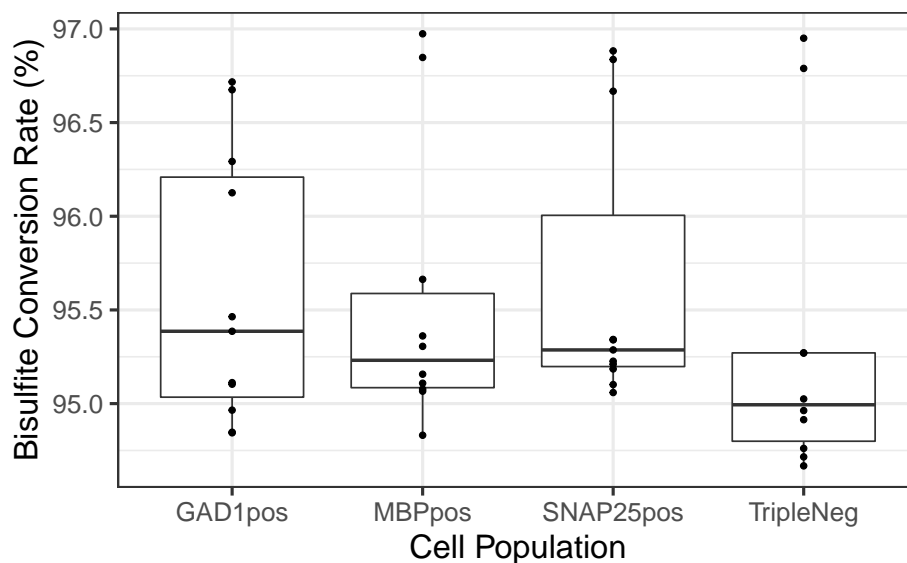


Figure 1: Bisulfite conversion rate by cell population across samples

Values center at 95% for all cell populations, ranging just over 2% total. The 'TripleNeg' population might appear to have lower conversion rates, but it's unclear if the difference is significant.

2.2 Relationship between Methylation Fractions across Cytosine Context, by Cell Population

Next, we'll explore if average methylation rate for each cytosine context correlates with that of other contexts across sample. For example, is a sample with highly methylated CpGs likely to have highly methylated CHGs? This is represented by the leftmost plot in middle row below.

```
# We'll make use of the 'ggpairs' function from the 'GGally' package, which is
# well-suited for comparison of various metrics against each other, providing
# density plots, comparison scatter plots, and correlation information.
ggpairs(
  meta_df,
  columns = c("perc_M_CpG", "perc_M_CHG", "perc_M_CHH"),
  xlab = "Methylation Rate (%)", ylab = "Methylation Rate (%)",
  columnLabels = c("CpG context", "CHG context", "CHH context"),
  mapping = aes(color = Cell.Population)
) +
  theme_bw(base_size = 15)
```

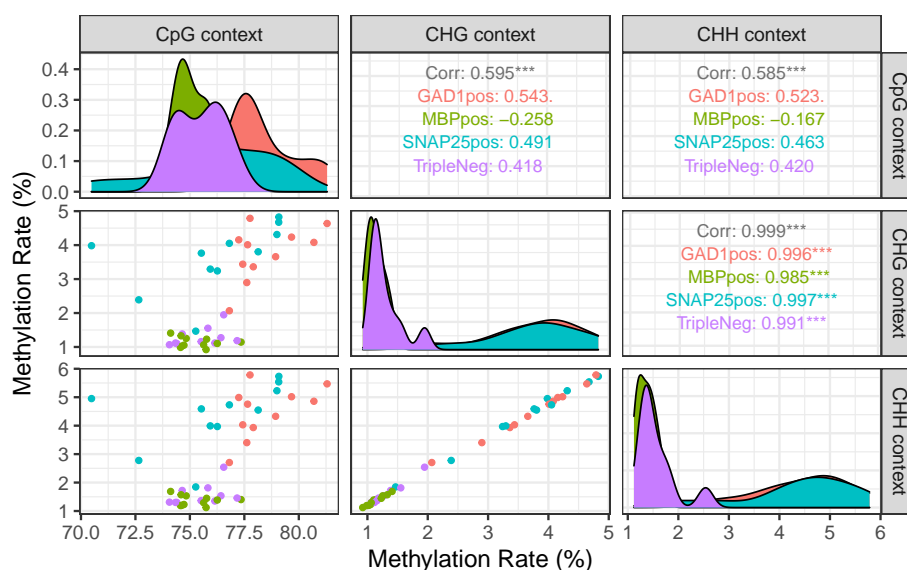


Figure 2: Comparison of average methylation rate by cytosine context and cell population

There is a correlation between average methylation rates of different cytosine contexts by sample. This is highly pronounced between CpH contexts (CHG vs. CHH). In each case, the relation appears roughly linear, though this is more questionable for the CpG vs. CHH context comparison. Generally, samples tend to cluster fairly well by cell population. Finally, for comparisons of CpG vs. CpH context, the strength of correlation between methylation rates varies significantly by cell type, with 'MBPpos' showing only weak correlation at best.

2.3 Fraction of Covered Cytosines by Cell Population

Another useful piece of information is how well-covered the genome is with methylation information. Does coverage of cytosines vary by a sample's cell type?

Example Analysis using BiocMAP Output Objects

```
# The matrices in 'assays(bs_cpg)' and 'assays(bs_cph)' are stored on disk. To
# speed up some below computations, we raise the per-block memory size
setAutoBlockSize(1e9)
## automatic block size set to 1e+09 bytes (was 1e+08)

# Get the proportion of cytosines in each object (context) that have non-zero
# coverage in at least one sample
meta_df$frac_covered_c_cpg <- DelayedArray::colMeans(assays(bs_cpg)$Cov > 3)
meta_df$frac_covered_c_cph <- DelayedArray::colMeans(assays(bs_cph)$Cov > 3)

# Convert our data frame to "long" format, with a CpG-associated and
# CpH-associated value for each coverage proportion
meta_df_long <- meta_df %>%
  pivot_longer(
    cols = starts_with("frac_covered_c"),
    names_to = "context",
    names_prefix = "frac_covered_c_",
    values_to = "frac_covered_c"
  )

# Plot fraction of covered cytosines by cell population
ggplot(
  meta_df_long,
  aes(x = Cell.Population, y = frac_covered_c, color = Cell.Population)
) +
  geom_boxplot(outlier.shape = NA) +
  geom_point(position = "jitter") +
  facet_wrap(~context) +
  labs(
    title = "Proportion of Covered Cytosines by Cytosine Context and Cell Population",
    x = "Cell Population",
    y = "Proportion of Covered Cytosines"
  ) +
  theme_bw(base_size = 10)
```

2.4 Distribution of Methylation Fractions across Cytosines by Cell Population

Grouping together all samples of a particular cell type, we'll explore the methylation-fraction distribution across cytosines, separately for both CpG and CpH contexts.

For the CpG cytosine context, we observe a bimodal distribution with peaks at fractions of 0 and 1. This suggests that within a particular sample, a CpG site is disproportionately likely to have consistent methylation pattern. For example, many CpG sites are such that all observations of the cytosine are methylated for a particular sample. Similarly, we don't see many sites where around half of the observed site are methylated for a given sample. It's also worth noting that the apparent bimodal form is likely not an artifact of low coverage— i.e., only a small fraction of sites are covered just once or twice, a circumstance that would cause over-representation of the fractions of 0 or 1.

Example Analysis using BiocMAP Output Objects

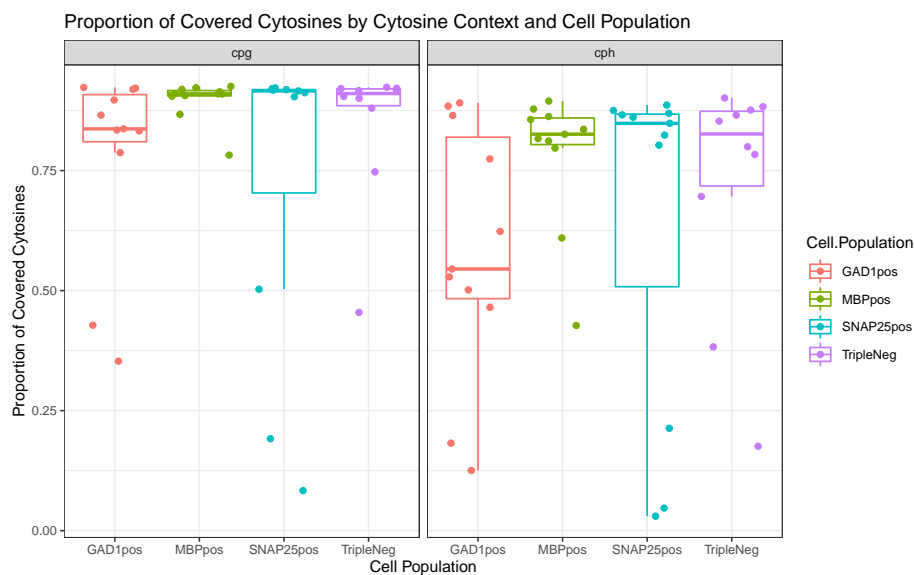


Figure 3: Proportion of covered cytosines by context and cell population

While around 80% of cytosines for a given sample have at least 3 observations in the experiment, there is wide variation between samples. At least 5 samples in each cytosine context have less than 50% sites 'well-covered', and CpH-context cytosines may have slightly less coverage generally than CpG-context cytosines.

```
# We'll use random sampling for computational efficiency. Set the seed for
# reproducibility
set.seed(444)

# Randomly subset to a particular number of cytosines, to both control memory
# and speed up plotting
max_sites <- 1000

# Look at CpG sites first
indices <- sample(nrow(bs_cpg), max_sites)
m_frac <- assays(bs_cpg)$M[indices, ] / assays(bs_cpg)$Cov[indices, ]

# It's worth looking at the distribution of coverage by site, since in theory
# this could be cause for the bimodality observed in the CpG portion of the
# plot below
head(table(assays(bs_cpg)$Cov[indices, ]), n = 20)
##
##      0      1      2      3      4      5      6      7      8      9     10     11     12     13
## 3597 1586 1486 1486 1499 1547 1659 1703 1726 1832 1953 1986 2071 1985
##      14     15     16     17     18     19
## 2032 1918 1834 1690 1560 1400

# Form a data frame for easy plotting: we'll collapse methylation data for all
# samples into a single column, 'm_frac'. Here 'LIBD.' denotes sample ID
meth_df <- data.frame(
  "m_frac" = as.numeric(m_frac),
```

Example Analysis using BiocMAP Output Objects

```
"LIBD." = rep(colnames(m_frac), each = max_sites),
"context" = "CpG"
)

# Now look at CpH sites
indices <- sample(nrow(bs_cph), max_sites)
m_frac <- assays(bs_cph)$M[indices, ] / assays(bs_cph)$Cov[indices, ]

# Again, we'll look at the distribution of coverage by site, to make sure it
# isn't leading to noisy methylation results
head(table(assays(bs_cph)$Cov[indices, ]), n = 20)
##
##      0      1      2      3      4      5      6      7      8      9     10     11     12     13
## 5010 2598 2806 3130 3302 3584 3644 3560 3283 2891 2475 1961 1462 1107
##      14      15      16      17      18      19
##      766      516      310      209      149      83

# Form a data frame for easy plotting: we'll collapse methylation data for all
# samples into a single column, 'm_frac'. Here 'LIBD.' denotes sample ID. The
# data frame will be appended to the existing CpG-associated one so we have
# data for both cytosine contexts together.
meth_df <- rbind(
  meth_df,
  data.frame(
    "m_frac" = as.numeric(m_frac),
    "LIBD." = rep(colnames(m_frac), each = max_sites),
    "context" = "CpH"
  )
)

# Label each observation (methylation fraction for a particular cytosine) with
# the cell population of the associated sample
meth_df$Cell.Population <- meta_df$Cell.Population[
  match(meth_df$"LIBD.", meta_df$"LIBD.")
]

ggplot(meth_df, aes(x = Cell.Population, y = m_frac, color = Cell.Population)) +
  geom_violin() +
  facet_wrap(~context) +
  labs(
    title = "Methylation Proportion across Cytosines by Cytosine Context and Cell Population",
    x = "Cell Population",
    y = "Methylation Proportion"
  ) +
  theme_bw(base_size = 10)
```

Next, we'll compute differentially methylated regions (DMRs) between cell populations. In particular, we'll compare MBP-positive samples to all other cell types, and plot methylation curves for CpG loci around the MBP genomic range.

Example Analysis using BiocMAP Output Objects

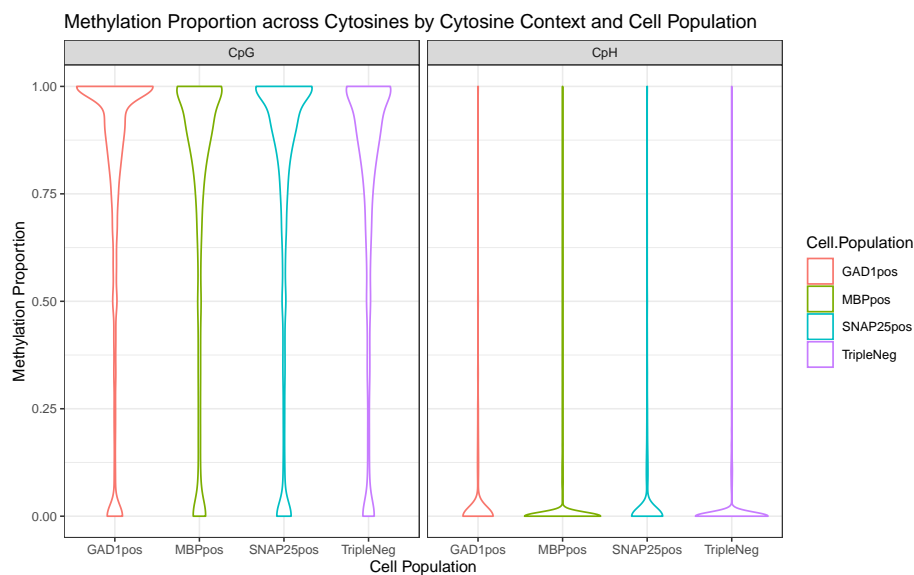


Figure 4: Methylation proportion across cytosines by context and cell population

For the CpG cytosine context, we observe a bimodal distribution with peaks at fractions of 0 and 1. This suggests that within a particular sample, a CpG site is disproportionately likely to have consistent methylation pattern. CpH sites are far less methylated than CpG sites. This reflects the well-known fact about DNA methylation that most methylation occurs at CpGs, though in some cells (such as neurons), CpH methylation plays an important role as well.

```
BiocParallel::register(BiocParallel::MulticoreParam(1))

cell_pop <- "MBPpos"

# To avoid false positive DMRs, we'll subset to CpGs where at least 5 samples
# in each group have at least two observations of the given CpG
num_cov_samples <- 5
num_cov_count <- 2

# Subset object based on coverage requirements
bs_cov <- getCoverage(bs_cpg)
loci_to_keep <- which(
  rowSums(
    bs_cov[, bs_cpg$Cell.Population == cell_pop] >= num_cov_count
  ) >= num_cov_samples &
  rowSums(
    bs_cov[, bs_cpg$Cell.Population != cell_pop] >= num_cov_count
  ) >= num_cov_samples
)
length(loci_to_keep)
## [1] 27691988

bs_cpg_sub <- bs_cpg[loci_to_keep, ]

# Define the two cell-population-based groups used to compute t-statistics
group1 <- unique(bs_cpg_sub$LIBD.[bs_cpg_sub$Cell.Population == cell_pop])
```


Example Analysis using BiocMAP Output Objects

```
group2 <- unique(bs_cpg_sub$"LIBD."[bs_cpg_sub$Cell.Population != cell_pop])

# Compute the t-stat and show the marginal distribution
cpg_t_stat <- BSmooth.tstat(
  bs_cpg_sub,
  group1 = group1, group2 = group2,
  estimate.var = "group2", local.correct = TRUE, verbose = TRUE
)
## [BSmooth.tstat] preprocessing ... done in 20.0 sec
## [BSmooth.tstat] computing stats within groups ... done in 246.7 sec
## [BSmooth.tstat] computing stats across groups ... done in 289.7 sec
```

```
plot(cpg_t_stat)
```

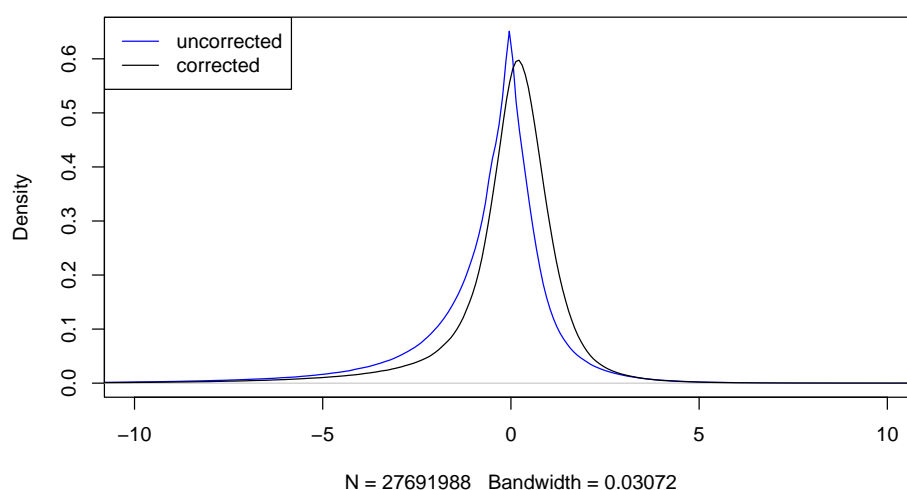


Figure 5: Marginal distribution of t-statistics, before and after correction

[The bsseq vignette](http://bioconductor.org/packages/release/bioc/vignettes/bsseq/inst/doc/bsseq_analysis.html#3_Computing_t-statistics) recommends local correction to improve the symmetry of the marginal distribution, an effect that can be seen here.

```
# Grab the genomic range associated with this particular cell population
if (cell_pop == "SNAP25pos") {
  gene_range <- GRanges("chr20:10172395-10308258")
} else if (cell_pop == "MBPpos") {
  gene_range <- GRanges("chr18:76978827-77133708")
} else if (cell_pop == "GAD1pos") {
  gene_range <- GRanges("chr2:170813210-170861151")
} else {
  stop(
    paste0("No gene associated with this cell population '", cell_pop, "'.")
  )
}

t_stat_df <- getStats(cpg_t_stat)

# Filters for the upcoming DMRs: minimum magnitude of t-stat and min number of
```

Example Analysis using BiocMAP Output Objects

```
# base pairs, respectively
thres <- sd(t_stat_df[, "tstat.corrected"]) * 2.5
n_bases <- 3
abs_mean_diff <- 0.1

# Compute DMRs and apply above filters
dmrs_orig <- dmrFinder(cpg_t_stat, cutoff = c(-1 * thres, thres))
## [dmrFinder] creating dmr data.frame
dmrs <- subset(dmr_orig, n >= n_bases & abs(meanDiff) >= abs_mean_diff)

# Color plots by cell population group
p_data <- pData(bs_cpg_sub)
p_data$col <- ifelse(bs_cpg_sub$Cell.Population == cell_pop, "red", "blue")
pData(bs_cpg_sub) <- p_data

# Plot the region around the gene associated with this particular cell
# population
plotRegion(
  bs_cpg_sub, gene_range,
  extend = width(gene_range) %/% 2, addRegions = dmrs
)
```

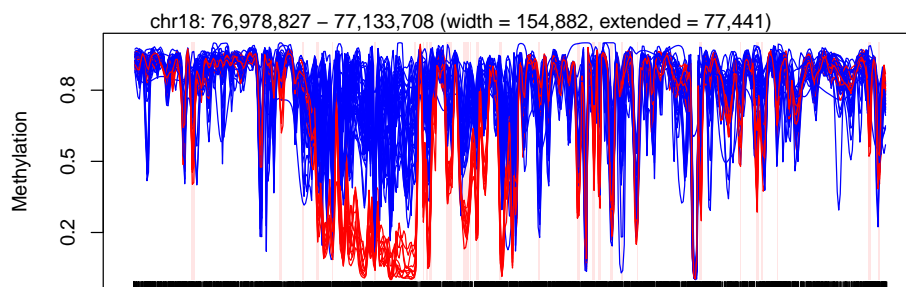


Figure 6: Methylation curves for MBP-positive samples (red) against other cell types (blue) around the MBP gene

To the eye, it appears that a fairly large region at the start of the MBP gene has differential methylation between groups; however, the only significant DMRs (highlighted in light red) are much narrower bands dispersed throughout the gene (and somewhat less so outside of the gene).

Date the vignette was generated:

```
## [1] "2022-03-01 13:37:10 EST"
```

Wall-clock time spent generating the vignette:

```
## Time difference of 1.461 hours
```

Memory-related information while generating the vignette:

```
gc()
##           used (Mb) gc trigger      (Mb)    max used     (Mb)
## Ncells   9417686   503   16825241   898.6   16825241   898.6
## Vcells 3144536751 23991 5883016828 44883.9 5883016828 44883.9
```

R session information:

Example Analysis using BiocMAP Output Objects

```
## - Session info -----
## setting value
## version R version 4.1.2 Patched (2021-11-04 r81138)
## os      CentOS Linux 7 (Core)
## system  x86_64, linux-gnu
## ui      X11
## language (EN)
## collate en_US.UTF-8
## ctype   en_US.UTF-8
## tz      US/Eastern
## date    2022-03-01
## pandoc  2.13 @ /jhpce/shared/jhpce/core/conda/miniconda3-4.6.14/envs/svnR-4.1.x/bin/ (via rmarkdown)
##
## - Packages -----
## package      * version date (UTC) lib source
## assertthat    0.2.1   2019-03-21 [2] CRAN (R 4.1.0)
## Biobase        * 2.54.0   2021-10-26 [2] Bioconductor
## BiocGenerics   * 0.40.0   2021-10-26 [2] Bioconductor
## BiocIO         1.4.0    2021-10-26 [2] Bioconductor
## BiocManager    1.30.16   2021-06-15 [2] CRAN (R 4.1.2)
## BiocParallel   1.28.3    2021-12-09 [2] Bioconductor
## BiocStyle      * 2.22.0   2021-10-26 [1] Bioconductor
## Biostrings     2.62.0    2021-10-26 [2] Bioconductor
## bitops         1.0-7     2021-04-24 [2] CRAN (R 4.1.0)
## bookdown       0.24      2021-09-02 [1] CRAN (R 4.1.2)
## BSgenome       1.62.0    2021-10-26 [2] Bioconductor
## bsseq          * 1.30.0   2021-10-26 [2] Bioconductor
## cli            3.2.0     2022-02-14 [2] CRAN (R 4.1.2)
## colorspace     2.0-3     2022-02-21 [2] CRAN (R 4.1.2)
## crayon         1.5.0     2022-02-14 [2] CRAN (R 4.1.2)
## data.table     1.14.2    2021-09-27 [2] CRAN (R 4.1.2)
## DBI            1.1.2     2021-12-20 [2] CRAN (R 4.1.2)
## DelayedArray   * 0.20.0   2021-10-26 [2] Bioconductor
## DelayedMatrixStats 1.16.0    2021-10-26 [2] Bioconductor
## digest         0.6.29    2021-12-01 [2] CRAN (R 4.1.2)
## dplyr          1.0.8     2022-02-08 [2] CRAN (R 4.1.2)
## ellipsis       0.3.2     2021-04-29 [2] CRAN (R 4.1.0)
## evaluate       0.15      2022-02-18 [2] CRAN (R 4.1.2)
## fansi          1.0.2     2022-01-14 [2] CRAN (R 4.1.2)
## farver         2.1.0     2021-02-28 [2] CRAN (R 4.1.0)
## fastmap        1.1.0     2021-01-25 [2] CRAN (R 4.1.0)
## generics       0.1.2     2022-01-31 [2] CRAN (R 4.1.2)
## GenomeInfoDb   * 1.30.1   2022-01-30 [2] Bioconductor
## GenomeInfoDbData 1.2.7     2021-11-01 [2] Bioconductor
## GenomicAlignments 1.30.0    2021-10-26 [2] Bioconductor
## GenomicRanges  * 1.46.1   2021-11-18 [2] Bioconductor
## GGally         * 2.1.2    2021-06-21 [2] CRAN (R 4.1.2)
## ggplot2        * 3.3.5    2021-06-25 [2] CRAN (R 4.1.0)
## glue           1.6.2     2022-02-24 [2] CRAN (R 4.1.2)
## gtable         0.3.0     2019-03-25 [2] CRAN (R 4.1.0)
## gtools         3.9.2     2021-06-06 [2] CRAN (R 4.1.0)
```

Example Analysis using BiocMAP Output Objects

```
## HDF5Array          * 1.22.1 2021-11-14 [2] Bioconductor
## htmltools          0.5.2 2021-08-25 [2] CRAN (R 4.1.2)
## httr               1.4.2 2020-07-20 [2] CRAN (R 4.1.0)
## IRanges            * 2.28.0 2021-10-26 [2] Bioconductor
## jsonlite           1.8.0 2022-02-22 [2] CRAN (R 4.1.2)
## knitr              1.37 2021-12-16 [2] CRAN (R 4.1.2)
## labeling           0.4.2 2020-10-20 [2] CRAN (R 4.1.0)
## lattice            0.20-45 2021-09-22 [3] CRAN (R 4.1.2)
## lifecycle          1.0.1 2021-09-24 [2] CRAN (R 4.1.2)
## limma              3.50.1 2022-02-17 [2] Bioconductor
## locfit             1.5-9.4 2020-03-25 [2] CRAN (R 4.1.0)
## lubridate          1.8.0 2021-10-07 [2] CRAN (R 4.1.2)
## magrittr           2.0.2 2022-01-26 [2] CRAN (R 4.1.2)
## Matrix             * 1.4-0 2021-12-08 [3] CRAN (R 4.1.2)
## MatrixGenerics     * 1.6.0 2021-10-26 [2] Bioconductor
## matrixStats        * 0.61.0 2021-09-17 [2] CRAN (R 4.1.2)
## munsell            0.5.0 2018-06-12 [2] CRAN (R 4.1.0)
## permute            0.9-7 2022-01-27 [2] CRAN (R 4.1.2)
## pillar             1.7.0 2022-02-01 [2] CRAN (R 4.1.2)
## pkgconfig          2.0.3 2019-09-22 [2] CRAN (R 4.1.0)
## plyr               1.8.6 2020-03-03 [2] CRAN (R 4.1.0)
## purrr              0.3.4 2020-04-17 [2] CRAN (R 4.1.0)
## R.methodsS3        1.8.1 2020-08-26 [2] CRAN (R 4.1.0)
## R.oo               1.24.0 2020-08-26 [2] CRAN (R 4.1.0)
## R.utils            2.11.0 2021-09-26 [2] CRAN (R 4.1.2)
## R6                 2.5.1 2021-08-19 [2] CRAN (R 4.1.2)
## RColorBrewer       1.1-2 2014-12-07 [2] CRAN (R 4.1.0)
## Rcpp               1.0.8 2022-01-13 [2] CRAN (R 4.1.2)
## RCurl              1.98-1.6 2022-02-08 [2] CRAN (R 4.1.2)
## RefManagerR        * 1.3.0 2020-11-13 [1] CRAN (R 4.1.2)
## reshape            0.8.8 2018-10-23 [2] CRAN (R 4.1.0)
## restfulr           0.0.13 2017-08-06 [2] CRAN (R 4.1.0)
## rhdf5              * 2.38.0 2021-10-26 [2] Bioconductor
## rhdf5filters        1.6.0 2021-10-26 [2] Bioconductor
## Rhdf5lib           1.16.0 2021-10-26 [2] Bioconductor
## rjson              0.2.21 2022-01-09 [2] CRAN (R 4.1.2)
## rlang              1.0.1 2022-02-03 [2] CRAN (R 4.1.2)
## rmarkdown          2.11 2021-09-14 [2] CRAN (R 4.1.2)
## Rsamtools          2.10.0 2021-10-26 [2] Bioconductor
## rtracklayer        1.54.0 2021-10-26 [2] Bioconductor
## S4Vectors          * 0.32.3 2021-11-21 [2] Bioconductor
## scales             1.1.1 2020-05-11 [2] CRAN (R 4.1.0)
## sessioninfo        * 1.2.2 2021-12-06 [2] CRAN (R 4.1.2)
## sparseMatrixStats  1.6.0 2021-10-26 [2] Bioconductor
## stringi            1.7.6 2021-11-29 [2] CRAN (R 4.1.2)
## stringr            1.4.0 2019-02-10 [2] CRAN (R 4.1.0)
## SummarizedExperiment * 1.24.0 2021-10-26 [2] Bioconductor
## tibble             3.1.6 2021-11-07 [2] CRAN (R 4.1.2)
## tidyr              * 1.2.0 2022-02-01 [2] CRAN (R 4.1.2)
## tidyselect         1.1.2 2022-02-21 [2] CRAN (R 4.1.2)
## utf8               1.2.2 2021-07-24 [2] CRAN (R 4.1.0)
```

Example Analysis using BiocMAP Output Objects

```
## vctrs          0.3.8    2021-04-29 [2] CRAN (R 4.1.0)
## withr          2.4.3    2021-11-30 [2] CRAN (R 4.1.2)
## xfun           0.29     2021-12-14 [2] CRAN (R 4.1.2)
## XML            3.99-0.9 2022-02-24 [2] CRAN (R 4.1.2)
## xml2           1.3.3    2021-11-30 [2] CRAN (R 4.1.2)
## XVector        0.34.0    2021-10-26 [2] Bioconductor
## yaml           2.3.5     2022-02-21 [2] CRAN (R 4.1.2)
## zlibbioc       1.40.0    2021-10-26 [2] Bioconductor
##
## [1] /users/neagles/R/4.1.x
## [2] /jhpce/shared/jhpce/core/conda/miniconda3-4.6.14/envs/svnR-4.1.x/R/4.1.x/lib64/R/site-library
## [3] /jhpce/shared/jhpce/core/conda/miniconda3-4.6.14/envs/svnR-4.1.x/R/4.1.x/lib64/R/library
##
## -----
```

3 Bibliography

This vignette was made possible by the following packages and software:

- R (R Core Team, 2021)
- *bsseq* = (Hansen, Langmead, and Irizarry, 2012)
- *BiocStyle* = (Oleś, 2021)
- *DelayedArray* = (Pagès, Hickey, and Lun, 2021)
- *GGally* = (Schloerke, Cook, Larmarange, Briatte, Marbach, Thoen, Elberg, and Crowley, 2021)
- *ggplot2* = (Wickham, 2016)
- *HDF5Array* = (Pagès, 2021)
- *knitr* = (Xie, 2021)
- *RefManageR* = (McLean, 2017)
- *rmarkdown* = (Allaire, Xie, McPherson, Luraschi, Ushey, Atkins, Wickham, Cheng, Chang, and Iannone, 2021)
- *sessioninfo* = (Wickham, Chang, Flight, Müller, and Hester, 2021)
- *tidyr* = (Wickham and Girlich, 2022)

[1] J. Allaire, Y. Xie, J. McPherson, et al. *rmarkdown*: Dynamic Documents for R. R package version 2.11. 2021. URL: <https://github.com/rstudio/rmarkdown>.

[2] K. D. Hansen, B. Langmead, and R. A. Irizarry. “BSmooth: from whole genome bisulfite sequencing reads to differentially methylated regions”. In: *Genome Biology* 13.10 (2012), p. R83. DOI: 10.1186/gb-2012-13-10-r83.

[3] M. W. McLean. “RefManageR: Import and Manage BibTeX and BibLaTeX References in R”. In: *The Journal of Open Source Software* (2017). DOI: 10.21105/joss.00338.

[4] A. Oleś. *BiocStyle*: Standard styles for vignettes and other Bioconductor documents. R package version 2.22.0. 2021. URL: <https://github.com/Bioconductor/BiocStyle>.

[5] H. Pagès. *HDF5Array*: HDF5 backend for DelayedArray objects. R package version 1.22.1. 2021. URL: <https://bioconductor.org/packages/HDF5Array>.

Example Analysis using BiocMAP Output Objects

- [6] H. Pagès, w. c. f. P. Hickey, and A. Lun. DelayedArray: A unified framework for working transparently with on-disk and in-memory array-like datasets. R package version 0.20.0. 2021. URL: <https://bioconductor.org/packages/DelayedArray>.
- [7] R Core Team. R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing. Vienna, Austria, 2021. URL: <https://www.R-project.org/>.
- [8] B. Schloerke, D. Cook, J. Larmarange, et al. GGally: Extension to 'ggplot2'. R package version 2.1.2. 2021. URL: <https://CRAN.R-project.org/package=GGally>.
- [9] H. Wickham. ggplot2: Elegant Graphics for Data Analysis. Springer-Verlag New York, 2016. ISBN: 978-3-319-24277-4. URL: <https://ggplot2.tidyverse.org>.
- [10] H. Wickham, W. Chang, R. Flight, et al. sessioninfo: R Session Information. R package version 1.2.2. 2021. URL: <https://CRAN.R-project.org/package=sessioninfo>.
- [11] H. Wickham and M. Girlich. tidyr: Tidy Messy Data. R package version 1.2.0. 2022. URL: <https://CRAN.R-project.org/package=tidyr>.
- [12] Y. Xie. knitr: A General-Purpose Package for Dynamic Report Generation in R. R package version 1.37. 2021. URL: <https://yihui.org/knitr/>.